

ANUBHAV SANKET

Bengaluru, India

📞 +91 8789378076 ✉ anubhavsanket22@gmail.com 🔗 [linkedin.com/in/anubhavsanket](https://www.linkedin.com/in/anubhavsanket) 🐙 github.com/anubhavsanket
🏠 kaggle.com/anubhavsanket 🌐 anubhavsanket.github.io/portfolio

Professional Summary

AI/ML Engineer building AI products end-to-end, from local LLM inference and RAG pipelines to deployed full-stack applications. Built privacy-first, offline-capable systems using Ollama, ChromaDB, LangGraph, FastAPI, React/TypeScript, Docker, WhisperX, FFmpeg, and CUDA. Strong focus on applied AI, developer tools, automation, model integration, and shipping production-oriented features for early-stage teams.

Skills

AI Engineering: LLMs, RAG, Agentic AI, LangGraph, Prompt Engineering, Embeddings, Vector Search, LLM Evaluation, Ollama, ChromaDB, multi-provider LLM routing, prompt caching, human-in-the-loop systems

Machine Learning: PyTorch, Transformers, scikit-learn, NLP, Computer Vision, WhisperX, OpenAI Whisper

Backend / Systems: Python, FastAPI, REST APIs, PostgreSQL, MongoDB, Redis, SQL

Infrastructure: Docker, Docker Compose, CUDA, NVIDIA GPU Inference, NVENC, Git, GitHub

Frontend: React, TypeScript, JavaScript, PyQt6, Tauri, Gradio, Vite, Tailwind CSS

Media / Multimodal: FFmpeg, WhisperX, Audio Processing, Video Processing, Word-level Transcription, Subtitle Generation, GPU-accelerated Media Pipelines

Analytics: Power BI, DAX, Excel, Pandas, NumPy, Portfolio Optimization

Professional Experience

Oaz Shakti

Nov. 2025 – Mar. 2026

Machine Learning / Power BI Intern

Remote

- Built an electrical-business analytics dashboard over 24K+ transactions worth 124.6M, tracking net revenue, gross profit of 23.49M, and an 18.9% margin to support operational and financial decision-making.
- Identified 8.4M in overdue payments through receivables analysis, surfacing a concrete cash-flow intervention area for the business.
- Developed and evaluated a quantitative asset-allocation solution for the Hull Tactical Market Prediction competition using Powell optimization, modified Sharpe-ratio maximization, volatility penalties, leverage bounds, and risk constraints.
- Achieved Kaggle Rank 158/4,506, placing in the global Top 3.5%; packaged the solution and supporting analysis in a public GitHub repository.

Varcons Technologies Pvt. Ltd.

Aug. 2023 – Oct. 2023

Full-stack Developer

Bengaluru, India

- Built and shipped a responsive ReactJS and MongoDB website for TechTune Motors with car-availability tracking, appointment scheduling, and real-time CRUD workflows.
- Improved page-load performance by 30% through database-query optimization and more efficient data access patterns.
- Worked across frontend implementation, backend integration, data modeling, and user-facing workflow design in a production-oriented web development setting.

Projects

Wayfarer: AI Job-Search Platform | *FastAPI, React/TypeScript, Redis, ChromaDB, Ollama, Docker*

- Built a full-stack AI platform that parses resumes, ranks live job posts, scores resume-job-description fit, and provides ATS-oriented feedback with human-in-the-loop review.
- Implemented Redis-backed caching and state management for LLM-powered job-matching workflows, supporting prompt reuse, rate-limit handling, and responsive processing of long-running tasks.
- Designed a retrieval pipeline combining in-process embeddings, vector search, hybrid relevance scoring, and graph-based resume memory to support token-efficient matching and explainable recommendations.
- Implemented a rate-limit-aware LLM router with NVIDIA NIM, OpenRouter, and Ollama providers, including fallback behavior and prompt caching for resilient inference across hosted and local models.
- Added configuration-driven connectors for 12 job boards and resume ingestion across 14 document formats using document parsing and OCR workflows.

- LocalBrainNotes: Local Agentic RAG Assistant** | *LangGraph, ChromaDB, LLM Evaluation, Ollama, Python*
- Built a local-first, Obsidian-like agentic RAG assistant over Markdown vaults, combining LangGraph orchestration, ChromaDB retrieval, and Ollama-based local inference.
 - Built a repeatable local LLM evaluation workflow using Ollama to compare agentic RAG responses across retrieval and prompting configurations, assessing relevance, grounding, consistency, and failure cases.
 - Focused the system on private knowledge access without requiring cloud-hosted model inference, demonstrating practical retrieval, agent orchestration, and local deployment trade-offs.
- Viral Clipper: Local GenAI Video Automation** | *Ollama, Gemma 3, WhisperX, PyQt6, FFmpeg, NVENC, CUDA*
- Developed an offline video-automation platform that uses LLM analysis to identify and extract high-value moments from long-form videos while keeping media processing local and privacy-preserving.
 - Implemented batch processing for 100+ parallel workflows, AI-assisted 9:16 reframing with face-centered composition, and animated word-level captions for short-form content generation.
 - Integrated WhisperX, FFmpeg, NVENC hardware encoding, PyQt6, and CUDA acceleration into a local desktop pipeline for transcription, clipping, reframing, caption rendering, and export.

Achievements & Publications

Kaggle Competition: Rank 158/4,506 (Top 3.5%) in the Hull Tactical - Market Prediction, 2025.

Research Publication: Customer Segmentation using K-Means clustering, published in IJNRD; applied Elbow Method and Silhouette Analysis to select clusters from purchasing and demographic data.

Google Tunix Hackathon Participant, 2025.

Core Member, Art & Literature Club; organized 3 events during Karnataka Rajyotsava 2022.

Education

Dayananda Sagar Academy of Technology and Management (VTU)	2020 – 2024
<i>B.E. in Artificial Intelligence and Machine Learning</i>	<i>Bengaluru, India</i>

Certifications

Microsoft Certified in Data Engineering on Azure

Power BI — Infosys Springboard

Cyber Security — Infosys Springboard

Hands-on Python — Coursera